



NoSQL Systems

Motivation

NoSQL: The Name

- “SQL” = Traditional relational DBMS
- Recognition over past decade or so:
Not every data management/analysis problem is best solved using a traditional relational DBMS
- “NoSQL” = “No SQL” =
Not using traditional relational DBMS
- “No SQL” \neq Don't use SQL language

NoSQL: The Name

- “SQL” = Traditional relational DBMS
- Recognition over past decade or so:
Not every data management/analysis problem is best solved *exclusively* using a traditional relational DBMS
- “NoSQL” = “No SQL” =
Not using traditional relational DBMS
- “No SQL” \neq Don’t use SQL language
- * “NoSQL” = “Not Only SQL”

Not every data management/analysis problem is best solved using a traditional DBMS

Database Management System (DBMS) provides....

... efficient, reliable, convenient, and safe
multi-user storage of and access to massive
amounts of persistent data.

Not every data management/analysis problem is best solved using a traditional DBMS

- Convenient
 - Multi-user
 - Safe
 - Persistent — files OK
 - Reliable — redo OK
 - Massive +++>
 - Efficient +++++>
- simple data model ✓
declarative query language ✓
transaction guarantees ✓

NoSQL Systems

Alternative to traditional relational DBMS

- + Flexible schema ✓
- + Quicker/cheaper to set up ✓
- + Massive scalability ✓
- + Relaxed consistency → higher performance & availability
- No declarative query language → more programming
- Relaxed consistency → fewer guarantees

Example #1: Web log analysis

Each record: UserID, URL, timestamp, additional-info ☆

Task: Load into database system

~~Data cleaning —~~

~~Data extraction —~~

~~Verification —~~

~~Schema specification —~~

Nothing

Example #1: Web log analysis

Each record: UserID, URL, timestamp, additional-info

Task: Find all records for...

- Given UserID ✓
- Given URL ✓
- Given timestamp ✓
- Certain construct appearing in additional-info

No SQL!

Highly
Parallelizable

Example #1: Web log analysis

Each record: UserID, URL, timestamp, additional-info

Task: Find all pairs of UserIDs accessing same URL

 Self-Join

Example #1: Web log analysis

Each record: UserID, URL, timestamp, additional-info

Separate records: UserID, name, age, gender, ...

Task: Find average age of user accessing given URL

SQL-like

Consistency

Example #2: Social-network graph

Each record: UserID₁, UserID₂

Separate records: UserID, name, age, gender, ...



Task: Find all friends of given user

Example #2: Social-network graph

Each record: UserID₁, UserID₂


Separate records: UserID, name, age, gender, ...

Task: Find all friends of friends given user



Example #2: Social-network graph

Each record: UserID₁, UserID₂ 

Separate records: UserID, name, age, gender, ... 

Task: Find all women friends of men friends of given user

Example #2: Social-network graph

Each record: UserID₁, UserID₂

Separate records: UserID, name, age, gender, ...

Task: Find all friends of friends of friends of ... friends of given user

Not suitable for SQL
X Consistency

Example #3: Wikipedia pages

Large collection of documents ✓

Combination of structured and unstructured data

Task: Retrieve introductory paragraph of all pages about U.S. presidents before 1900

Consistency

NoSQL Systems

Alternative to traditional relational DBMS

- + Flexible schema ✓
- + Quicker/cheaper to set up ✓
- + Massive scalability ✓
- + Relaxed consistency → higher performance & availability
- No declarative query language → more programming
- Relaxed consistency → fewer guarantees