




NoSQL Systems

Overview
(as of November 2011)

NoSQL Systems

- Not every data management/analysis problem is best solved *exclusively* using a traditional DBMS
- “NoSQL” = “Not Only SQL” 

NoSQL Systems

Alternative to traditional relational DBMS

- + Flexible schema ✓
- + Quicker/cheaper to set up ✓
- + Massive scalability ✓
- + Relaxed consistency ✓ → higher performance & availability
- No declarative query language → more programming
- Relaxed consistency → fewer guarantees

NoSQL Systems

Several incarnations

- MapReduce framework ~ OLAP
- Key-value stores ~ OLTP
- Document stores
- Graph database systems

Column Stores

MapReduce Framework

Originally from Google, open source Hadoop

- No data model, data stored in files

GFS
HDFS

- User provides specific functions

map() reduce()
reader() writer() combiner()

- System provides data processing “glue”, fault-tolerance, scalability

Map and Reduce Functions

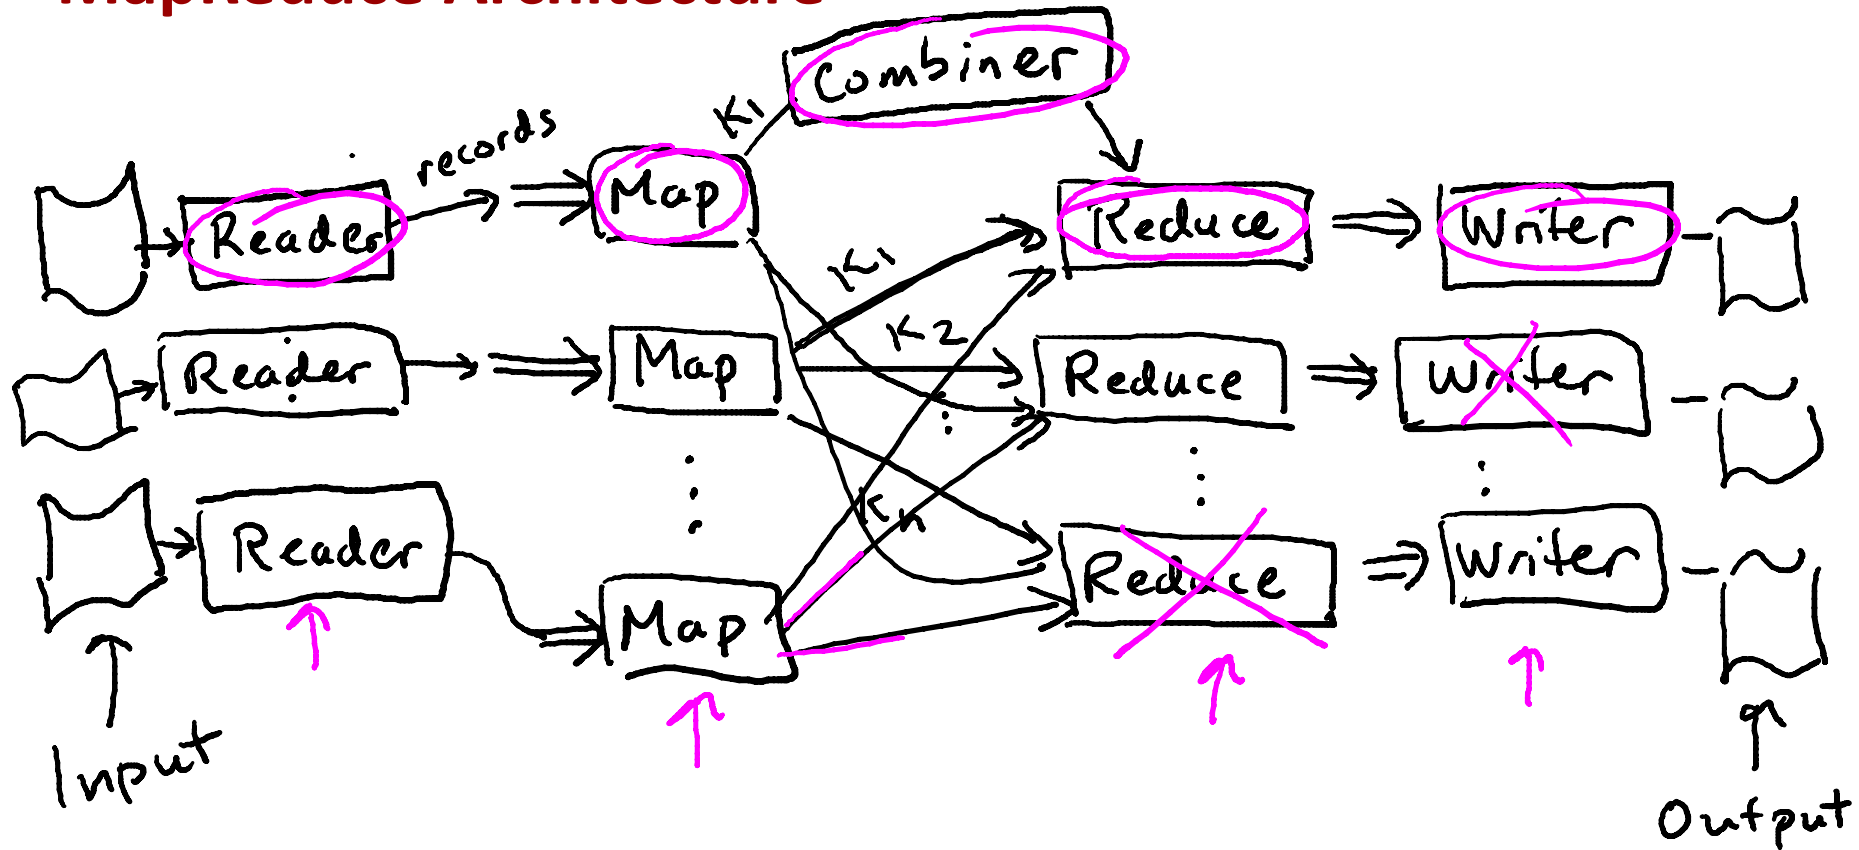
Map: Divide problem into subproblems

$\text{map}(\text{item}) \rightarrow 0 \text{ or more } \langle \text{key}, \text{value} \rangle \text{ pairs}$

Reduce: Do work on subproblems, combine results

$\text{reduce}(\text{key}, \text{list-of-values}) \rightarrow 0 \text{ or more records}$

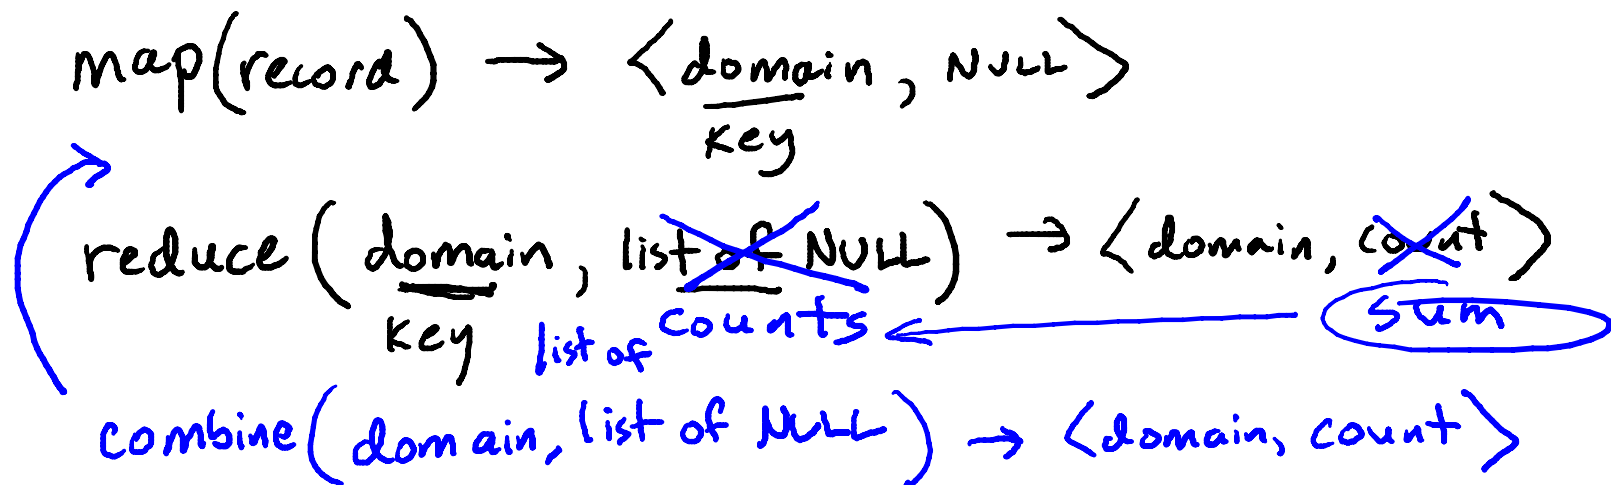
MapReduce Architecture



MapReduce Example: Web log analysis

Each record: UserID, URL, timestamp, additional-info ←

Task: Count number of accesses for each domain (inside URL)



MapReduce Example (modified #1)

Each record: UserID, URL, timestamp, additional-info

Task: Total "value" of accesses for each domain based on additional-info ←

map(record) → <domain, score>

reduce(domain, list of scores) → <domain, sum>

MapReduce Example (modified #2)

Each record: UserID, URL, timestamp, additional-info

Separate records: UserID, name, age, gender, ...



Task: Total “value” of accesses for each domain based on user attributes

MapReduce Framework

- No data model, data stored in files ✓
- User provides specific functions ✓
- System provides data processing “glue”, fault-tolerance, scalability ✓

MapReduce Framework

Schemas and declarative queries are missed

Hive – schemas, SQL-like query language

Pig – more imperative but with relational operators

- Both compile to “workflow” of Hadoop (MapReduce) jobs

Dryad allows user to specify workflow

- Also DryadLINQ language

Key-Value Stores “OLTP”

Extremely simple interface

- Data model: (key, value) pairs
- Operations: Insert(key, value), Fetch(key),
→ Update(key), Delete(key)

Implementation: efficiency, scalability, fault-tolerance

- Records distributed to nodes based on key
- Replication
- Single-record transactions, “eventual consistency”

Key-Value Stores

Extremely simple interface

- Data model: (key, value) pairs *structure*
- Operations: Insert(key, value), Fetch(key), Update(key), Delete(key)
- Some allow (non-uniform) columns within value
- Some allow Fetch on range of keys

Example systems

$$2 < \text{key} < 10$$

- Google BigTable, Amazon Dynamo, Cassandra, Voldemort, HBase, ...

Document Stores

Like Key-Value Stores except value is document


- Data model: (key, document) pairs ←
- Document: JSON, XML, other semistructured formats
- Basic operations: Insert(key,document), Fetch(key),
→ Update(key), Delete(key) ←
- Also Fetch based on document contents ← *System/format specific*

Example systems

- CouchDB, MongoDB, SimpleDB, ...

Graph Database Systems

- Data model: nodes and edges
- Nodes may have properties (including ID)
- Edges may have labels or roles



Graph Database Systems

- Interfaces and query languages vary
- Single-step versus “path expressions” versus full recursion
- Example systems
 - Neo4j, FlockDB, Pregel, ...
- RDF “triple stores” can map to graph databases

NoSQL Systems

- “NoSQL” = “Not Only SQL”

Not every data management/analysis problem is best solved *exclusively* using a traditional DBMS


- Current incarnations

- MapReduce framework ✓

- Key-value stores ✓

- Document stores ✓

- Graph database systems ✓



NoSQL Systems

Overview
(as of November 2011)